

Can I eat this? — Compass Group AI Engineer pre-work

12:40. Third in the queue.

Priya has a severe nut allergy. The menu is a laminated sheet. The allergen folder is behind the counter. The server started Monday.

She has about thirty seconds to decide if lunch is safe.

Seyi Bello · a two-tap allergen check for that exact moment

The problem

The data is the actual hazard, not the queue

Veggie chilli (vegan?)

"celery. may contain traces of nuts"

Thai green curry

"see chef"

Katsu sando (special)

"ask counter"

Fruit pot

"null"

This is the brief's own starter menu, not invented. A server who started Monday can't parse this in thirty seconds. Neither should an AI have to guess.

The design

Two taps. One of three honest answers.

Safe

Confirmed clean against the allergen asked about.

Unsafe

Confirmed to contain it.

Ask staff

Data is ambiguous or missing. Never guessed as safe.

Not a confidence score. A score just moves the guessing problem onto someone with thirty seconds, not the time to interpret "72% safe." "Ask staff" is the correct answer whenever the data itself doesn't support a claim of safe, and it can never be silently overridden.

Live demo

Now, on the actual app

Can I eat this?

Pick a dish, then pick an allergen. Two taps, no typing.

1. Select a dish

Chicken katsu curry £5.20	Veggie chilli (vegan?) £4.50	Margherita pizza slice £2.80	Thal green curry £5.20
Jacket potato with beans £3.95	Flapjack £1.60	Halloumi wrap £4.80	Beef lasagne £5.50
Fruit pot £2.10	Katsu sando (special) £4.95	Prawn cocktail £3.80	Fish finger sandwich £4.20
Moules marinière £6.50	Falafel wrap £4.60	Lentil dahl £4.10	Chicken satay skewers £5.10
Sausage roll £2.50	Egg fried rice £3.60	Lupin seed bread roll £1.20	Coleslaw (side) £1.80
Sticky toffee pudding £3.20	Vegetable spring rolls £2.90	Chickpea curry £4.30	Cheese & pickle sandwich £3.40
Smoked salmon bagel £5.80	Beef burger (no bun) £4.90	Tuna mayo baguette £4.10	Almond croissant £2.60

- Margherita pizza + milk → confirmed unsafe
- Veggie chilli + nuts → ask staff, not a guess
- Live at frontend-rose-three-32.vercel.app

Can I eat this?

Pick a dish, then pick an allergen. Two taps, no typing.

Can't confirm - ask staff

Veggie chilli (vegan?)

Menu data for "Veggie chilli (vegan?)" is flagged ambiguous (may contain traces of nuts) - cannot confirm safe.

Ask about another allergen

New dish

Built for the counter, not just the queue

A QR code on the counter opens this

[← full picker](#)

Nuts - whole menu

1 unsafe · 11 ask staff · 18 safe

Almond croissant

Almond croissant is confirmed to contain nuts.

Unsafe

Veggie chilli (vegan?)

Menu data for 'Veggie chilli (vegan?)' is flagged ambiguous (may contain traces of nuts) - cannot confirm safe.

Ask staff

Thai green curry

Menu data for 'Thai green curry' is flagged ambiguous (see chef - allergen info not on the board, must ask kitchen staff) - cannot confirm safe.

Ask staff

Flapjack

Menu data for 'Flapjack' is flagged ambiguous (peanuts?? - listed with question marks in original notes, presence not confirmed) - cannot confirm safe.

Ask staff

Fruit pot

Menu data for 'Fruit pot' is flagged ambiguous (no allergen data recorded (field literally reads 'null')) - cannot confirm safe.

Ask staff

Katsu sando (special)

Menu data for 'Katsu sando (special)' is flagged ambiguous (ask counter - no allergen data available, must ask staff directly) - cannot confirm safe.

Ask staff


Lentil dahl

Menu data for 'Lentil dahl' is flagged ambiguous (allergen field left blank on the board) - cannot confirm safe.

Ask staff

- The whole menu, scanned for one allergen
- Sorted worst-first, unsafe before ask-staff before safe
- No dish-picking, no typing
- Built for Priya, and for the server who doesn't know the menu yet either

Two routes. One touches an LLM.



Button taps skip the LLM entirely, a dict lookup under 200ms. GPT-4o-mini only runs for free-text questions.

The questions I expect to defend

Why these decisions, specifically

Why not a confidence score?

It moves the guessing problem onto someone with thirty seconds. Three plain states don't.

Why not agentic, why no LangGraph or multi-agent setup?

This is one bounded classification decision, not a multi-step workflow. Agentic complexity here is tool-stacking, not creativity.

Why no vector database?

30 structured records, not a large unstructured corpus. Fuzzy string matching solves it for free, with no risk of confidently returning the wrong dish.

What I'd build next

With another week

- 01** A batch pipeline to digitise a photographed menu into structured entries, hand-curating doesn't scale past a demo
- 02** Admin auth and a proper editing view for updating menu data day to day
- 03** The eval, observe, diagnose, release loop, once this is serving real traffic rather than a one-off demo

[Back to Priya](#)

Thirty seconds. Two taps. One honest answer.


Including "ask staff" when that's the truth.

Live demo	frontend-rose-three-32.vercel.app
API	can-i-eat-this-api.fly.dev
Code	github.com/seyiabello/can-i-eat-this

Questions?

Backup slide — only if asked about production

What's next for the ops loop



What's built today (trace log, state check) versus the full evaluate/observe/diagnose/release loop, deliberately not built yet, that's solving a maintenance problem this prototype doesn't have.